

Section 1: Getting Started With Pandas

We will begin by introducing the `Series`, `DataFrame`, and `Index` classes, which are the basic building blocks of the `pandas` library, and showing how to work with them. By the end of this section, you will be able to create `DataFrames` and perform operations on them to inspect and filter the data.

Anatomy of a DataFrame

A **DataFrame** is composed of one or more **Series**. The names of the **Series** form the column names, and the row labels form the **Index**.

In [1]:

Out[1]:

	name	id	nametype	recclass	mass (g)	fall	year	reclat	reclong	GeoL
0	Aachen	1	Valid	L5	21	Fell	01/01/1880 12:00:00 AM	50.77500	6.08333	(6
1	Aarhus	2	Valid	H6	720	Fell	01/01/1951 12:00:00 AM	56.18333	10.23333	(56 10
2	Abee	6	Valid	EH4	107000	Fell	01/01/1952 12:00:00 AM	54.21667	-113.00000	(52
3	Acapulco	10	Valid	Acapulcoite	1914	Fell	01/01/1976 12:00:00 AM	16.88333	-99.90000	(16
4	Achiras	370	Valid	L6	780	Fell	01/01/1902 12:00:00 AM	-33.16667	-64.95000	(-32

Source: [NASA's Open Data Portal](#)

Series:

In [2]:

```
Out[2]: 0      Aachen
        1      Aarhus
        2      Abee
        3  Acapulco
        4  Achiras
        Name: name, dtype: object
```

Columns:

```
In [3]:
```

```
Out[3]: Index(['name', 'id', 'nametype', 'recclass', 'mass (g)', 'fall', 'year',
              'reclat', 'reclong', 'GeoLocation'],
              dtype='object')
```

Index:

```
In [4]:
```

```
Out[4]: RangeIndex(start=0, stop=5, step=1)
```

Creating DataFrames

We can create DataFrames from a variety of sources such as other Python objects, flat files, webscraping, and API requests. Here, we will see just a couple of examples, but be sure to check out [this page](#) in the documentation for a complete list.

Using a flat file

```
In [5]:
```

Tip: There are many parameters to this function to handle some initial processing while reading in the file – be sure check out the [documentation](#).

Using data from an API

Collect the data from [NASA's Open Data Portal](#) using the Socrata Open Data API (SODA) with the `requests` library:

```
In [6]:
```

Create the DataFrame with the resulting payload:

In [7]:

Out[7]:

	name	id	nametype	recclass	mass	fall	year	reclat	reclong	geoloc
0	Aachen	1	Valid	L5	21	Fell	1880-01-01T00:00:00.000	50.77500	6.08333	{'lat': 50.775, 'long': 6.08333}
1	Aarhus	2	Valid	H6	720	Fell	1951-01-01T00:00:00.000	56.18333	10.23333	{'lat': 56.18333, 'long': 10.23333}
2	Abee	6	Valid	EH4	107000	Fell	1952-01-01T00:00:00.000	54.21667	-113.00000	{'lat': 54.21667, 'long': -113.0}

Tip: `df.to_csv('data.csv')` writes this data to a new file called `data.csv`.

Inspecting the data

Now that we have some data, we need to perform an initial inspection of it. This gives us information on what the data looks like, how many rows/columns there are, and how much data we have.

Let's inspect the `meteorites` data.

How many rows and columns are there?

In [8]:

Out[8]: (45716, 10)

What are the column names?

In [9]:

```
Out[9]: Index(['name', 'id', 'nametype', 'recclass', 'mass (g)', 'fall', 'year',
              'reclat', 'reclong', 'GeoLocation'],
              dtype='object')
```

What type of data does each column currently hold?

In [10]:

```
Out[10]: name          object
         id            int64
         nametype      object
         recclass      object
         mass (g)      float64
         fall          object
         year          object
         reclat        float64
         reclong       float64
         GeoLocation   object
         dtype: object
```

What does the data look like?

In [11]:

Out[11]:

	name	id	nametype	recclass	mass (g)	fall	year	reclat	reclong	Geo
0	Aachen	1	Valid	L5	21.0	Fell	01/01/1880 12:00:00 AM	50.77500	6.08333	
1	Aarhus	2	Valid	H6	720.0	Fell	01/01/1951 12:00:00 AM	56.18333	10.23333	(1
2	Abee	6	Valid	EH4	107000.0	Fell	01/01/1952 12:00:00 AM	54.21667	-113.00000	(
3	Acapulco	10	Valid	Acapulcoite	1914.0	Fell	01/01/1976 12:00:00 AM	16.88333	-99.90000	('
4	Achiras	370	Valid	L6	780.0	Fell	01/01/1902 12:00:00 AM	-33.16667	-64.95000	(-

Sometimes there may be extraneous data at the end of the file, so checking the bottom few rows is also important:

In [12]:

Out [12]:

	name	id	nametype	recclass	mass (g)	fall	year	reclat	reclo
45711	Zillah 002	31356	Valid	Eucrite	172.0	Found	01/01/1990 12:00:00 AM	29.03700	17.018
45712	Zinder	30409	Valid	Pallasite, ungrouped	46.0	Found	01/01/1999 12:00:00 AM	13.78333	8.966
45713	Zlin	30410	Valid	H4	3.3	Found	01/01/1939 12:00:00 AM	49.25000	17.666
45714	Zubkovsky	31357	Valid	L6	2167.0	Found	01/01/2003 12:00:00 AM	49.78917	41.504
45715	Zulu Queen	30414	Valid	L3.7	200.0	Found	01/01/1976 12:00:00 AM	33.98333	-115.683

Get some information about the DataFrame

In [13]:

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45716 entries, 0 to 45715
Data columns (total 10 columns):
#   Column          Non-Null Count  Dtype
---  -
0   name            45716 non-null  object
1   id              45716 non-null  int64
2   nametype        45716 non-null  object
3   recclass        45716 non-null  object
4   mass (g)        45585 non-null  float64
5   fall            45716 non-null  object
6   year            45425 non-null  object
7   reclat          38401 non-null  float64
8   reclong         38401 non-null  float64
9   GeoLocation     38401 non-null  object
dtypes: float64(3), int64(1), object(6)
memory usage: 3.5+ MB

```

Extracting subsets

A crucial part of working with DataFrames is extracting subsets of the data: finding rows that meet a certain set of criteria, isolating columns/rows of interest, etc. After narrowing down our data, we are closer to discovering insights. This section will be the backbone of many analysis tasks.

Selecting columns

We can select columns as attributes if their names would be valid Python variables:

In [14]:

Out[14]:

```
0          Aachen
1          Aarhus
2           Abee
3        Acapulco
4         Achiras
...
45711    Zillah 002
45712         Zinder
45713         Zlin
45714    Zubkovsky
45715    Zulu Queen
Name: name, Length: 45716, dtype: object
```

If they aren't, we have to select them as keys. However, we can select multiple columns at once this way:

In [15]:

Out [15]:

	name	mass (g)
0	Aachen	21.0
1	Aarhus	720.0
2	Abee	107000.0
3	Acapulco	1914.0
4	Achiras	780.0
...
45711	Zillah 002	172.0
45712	Zinder	46.0
45713	Zlin	3.3
45714	Zubkovsky	2167.0
45715	Zulu Queen	200.0

45716 rows × 2 columns

Selecting rows

In [16]:

Out [16]:

	name	id	nametype	recclass	mass (g)	fall	year	reclat	reclong
100	Benton	5026	Valid	LL6	2840.0	Fell	01/01/1949 12:00:00 AM	45.95000	-67.55000
101	Berduc	48975	Valid	L6	270.0	Fell	01/01/2008 12:00:00 AM	-31.91000	-58.32833
102	Béréba	5028	Valid	Eucrite-mmict	18000.0	Fell	01/01/1924 12:00:00 AM	11.65000	-3.65000
103	Berlanguillas	5029	Valid	L6	1440.0	Fell	01/01/1811 12:00:00 AM	41.68333	-3.80000

Indexing

We use `iloc[]` to select rows and columns by their position:

In [17]:

Out[17]:

	name	recclass	mass (g)	year
100	Benton	LL6	2840.0	01/01/1949 12:00:00 AM
101	Berduc	L6	270.0	01/01/2008 12:00:00 AM
102	Béréba	Eucrite-mmict	18000.0	01/01/1924 12:00:00 AM
103	Berlanguillas	L6	1440.0	01/01/1811 12:00:00 AM

We use `loc[]` to select by name:

In [18]:

Out[18]:

	mass (g)	fall	year
100	2840.0	Fell	01/01/1949 12:00:00 AM
101	270.0	Fell	01/01/2008 12:00:00 AM
102	18000.0	Fell	01/01/1924 12:00:00 AM
103	1440.0	Fell	01/01/1811 12:00:00 AM
104	960.0	Fell	01/01/2004 12:00:00 AM

Filtering with Boolean masks

A **Boolean mask** is a array-like structure of Boolean values – it's a way to specify which rows/columns we want to select (`True`) and which we don't (`False`).

Here's an example of a Boolean mask for meteorites weighing more than 50 grams that was found on Earth (i.e. they were not observed falling):

In [19]:


```
Out[19]: 0      False
         1      False
         2      False
         3      False
         4      False
         ...
        45711    True
        45712    False
        45713    False
        45714    True
        45715    True
        Length: 45716, dtype: bool
```

Important: Take note of the syntax here. We surround each condition with parentheses, and we use bitwise operators (`&` , `|` , `~`) instead of logical operators (`and` , `or` , `not`).

We can use this Boolean mask to select the full row for each of these meteorites:

```
In [20]:
```

Out [20]:

	name	id	nametype	recclass	mass (g)	fall	year	reclat	re
37	Northwest Africa 5815	50693	Valid	L5	256.80	Found	NaN	0.00000	0.0
757	Dominion Range 03239	32591	Valid	L6	69.50	Found	01/01/2002 12:00:00 AM	NaN	
804	Dominion Range 03240	32592	Valid	LL5	290.90	Found	01/01/2002 12:00:00 AM	NaN	
1111	Abajo	4	Valid	H5	331.00	Found	01/01/1982 12:00:00 AM	26.80000	-105.0
1112	Abar al' Uj 001	51399	Valid	H3.8	194.34	Found	01/01/2008 12:00:00 AM	22.72192	48.0
...
45709	Zhongxiang	30406	Valid	Iron	100000.00	Found	01/01/1981 12:00:00 AM	31.20000	112.0
45710	Zillah 001	31355	Valid	L6	1475.00	Found	01/01/1990 12:00:00 AM	29.03700	17.0
45711	Zillah 002	31356	Valid	Eucrite	172.00	Found	01/01/1990 12:00:00 AM	29.03700	17.0
45714	Zubkovsky	31357	Valid	L6	2167.00	Found	01/01/2003 12:00:00 AM	49.78917	41.0
45715	Zulu Queen	30414	Valid	L3.7	200.00	Found	01/01/1976 12:00:00 AM	33.98333	-115.0

18854 rows × 10 columns

*Tip: Boolean masks can be used with `loc[]` and `iloc[]`.*An alternative to this is the `query()` method:

In [21]:

Out[21]:

	name	id	nametype	recclass	mass (g)	fall	year	reclat	re
37	Northwest Africa 5815	50693	Valid	L5	256.80	Found	NaN	0.00000	0.0
757	Dominion Range 03239	32591	Valid	L6	69.50	Found	01/01/2002 12:00:00 AM	NaN	
804	Dominion Range 03240	32592	Valid	LL5	290.90	Found	01/01/2002 12:00:00 AM	NaN	
1111	Abajo	4	Valid	H5	331.00	Found	01/01/1982 12:00:00 AM	26.80000	-105.0
1112	Abar al' Uj 001	51399	Valid	H3.8	194.34	Found	01/01/2008 12:00:00 AM	22.72192	48.0
...
45709	Zhongxiang	30406	Valid	Iron	100000.00	Found	01/01/1981 12:00:00 AM	31.20000	112.0
45710	Zillah 001	31355	Valid	L6	1475.00	Found	01/01/1990 12:00:00 AM	29.03700	17.0
45711	Zillah 002	31356	Valid	Eucrite	172.00	Found	01/01/1990 12:00:00 AM	29.03700	17.0
45714	Zubkovsky	31357	Valid	L6	2167.00	Found	01/01/2003 12:00:00 AM	49.78917	41.0
45715	Zulu Queen	30414	Valid	L3.7	200.00	Found	01/01/1976 12:00:00 AM	33.98333	-115.0

18854 rows × 10 columns

Tip: Here, we can use both logical operators and bitwise operators.

Calculating summary statistics

In the next section of this workshop, we will discuss data cleaning for a more meaningful analysis of these datasets; however, we can already extract some interesting insights from the `meteorites` data by calculating summary statistics.

How many of the meteorites were found versus observed falling?

In [22]:

```
Out[22]: Found      44609
 Fell        1107
 Name: fall, dtype: int64
```

Tip: Pass in `normalize=True` to see this result as percentages. Check the [documentation](#) for additional functionality.

What was the mass of the average meteorite?

In [23]:

```
Out[23]: 32.6
```

We can take this a step further and look at quantiles:

In [24]:

```
Out[24]: 0.01      0.44
 0.05      1.10
 0.95    4000.00
 0.99   50600.00
 Name: mass (g), dtype: float64
```

What was the mass of the heaviest meteorite?

In [25]:

```
Out[25]: 60000000.0
```

Let's extract the information on this meteorite:

In [26]:

```
Out[26]: name                Hoba
 id                11890
 nametype          Valid
 recclass          Iron, IVB
 mass (g)          60000000.00000
 fall              Found
 year              01/01/1920 12:00:00 AM
 reclat            -19.58333
 reclong           17.91667
 GeoLocation      (-19.58333, 17.91667)
 Name: 16392, dtype: object
```

How many different types of meteorite classes are represented in this dataset?

In [27]:

Out[27]: 466

Some examples:

In [28]:

Out[28]: array(['L5', 'H6', 'EH4', 'Acapulcoite', 'L6', 'LL3-6', 'H5', 'L',
'Diogenite-pm', 'Unknown', 'H4', 'H', 'Iron, IVA', 'CR2-an', 'LL5',
'CI1', 'L/LL4', 'Eucrite-mmict', 'CV3', 'Ureilite-an'],
dtype=object)

Note: All fields preceded with "rec" are the values recommended by The Meteoritical Society. Check out [this Wikipedia article](#) for some information on meteorite classes.

Get some summary statistics on the data itself

We can get common summary statistics for all columns at once. By default, this will only be numeric columns, but here, we will summarize everything together:

In [29]:

Out[29]:

	name	id	nametype	recclass	mass (g)	fall	year	
count	45716	45716.000000	45716	45716	4.558500e+04	45716	45425	3840
unique	45716	NaN	2	466	NaN	2	266	
top	Northwest Africa 890	NaN	Valid	L6	NaN	Found	01/01/2003 12:00:00 AM	
freq	1	NaN	45641	8285	NaN	44609	3323	
mean	NaN	26889.735104	NaN	NaN	1.327808e+04	NaN	NaN	-3
std	NaN	16860.683030	NaN	NaN	5.749889e+05	NaN	NaN	4
min	NaN	1.000000	NaN	NaN	0.000000e+00	NaN	NaN	-8
25%	NaN	12688.750000	NaN	NaN	7.200000e+00	NaN	NaN	-7
50%	NaN	24261.500000	NaN	NaN	3.260000e+01	NaN	NaN	-7
75%	NaN	40656.750000	NaN	NaN	2.026000e+02	NaN	NaN	
max	NaN	57458.000000	NaN	NaN	6.000000e+07	NaN	NaN	8

Important: NaN values signify missing data. For instance, the `fall` column contains strings, so there is no value for `mean`; likewise, `mass (g)` is numeric, so we don't have entries for the categorical summary statistics (`unique`, `top`, `freq`).

Check out the documentation for more descriptive statistics:

- [Series](#)
- [DataFrame](#)

Up Next: Data Wrangling

Let's take a 10-minute break for some exercises to check your understanding:

1. Create a DataFrame by reading in the `2019_Yellow_Taxi_Trip_Data.csv` file.
2. Find the dimensions (number of rows and number of columns) in the data.
3. Calculate summary statistics for the `fare_amount`, `tip_amount`, `tolls_amount`, and `total_amount` columns.
4. Isolate the `fare_amount`, `tip_amount`, `tolls_amount`, and `total_amount` for the longest trip (`trip_distance`).

Exercises

1. Create a DataFrame by reading in the `2019_Yellow_Taxi_Trip_Data.csv` file:

In []:

2. Find the dimensions (number of rows and number of columns) in the data:

In []:

3. Calculate summary statistics for the `fare_amount`, `tip_amount`, `tolls_amount`, and `total_amount` columns:

In []:

4. Isolate the `fare_amount`, `tip_amount`, `tolls_amount`, and `total_amount` for the longest trip (`trip_distance`):

In []: